

Noise Removal versus Signal Loss in Network Filtering of Daily GNSS Time Series

H. Oğuz Bolat

Department of Geomatics Engineering, Yıldız Technical University, Istanbul, Türkiye
oguz.bolat@std.yildiz.edu.tr

Research proposal, October 2026

Abstract

Daily GNSS position time series are routinely filtered to suppress noise that is shared across stations, but the same filters can remove real transient deformation, because slow slip, afterslip and loading signals are spatially coherent too. Common-mode filters have been judged mainly by how much scatter they remove, a criterion that also rewards removing signal. Where signal preservation has been examined, it has been inferred from the consistency of noisy observed offsets or tested with synthetic signals in synthetic noise, and never as a function of the size of a transient relative to the reach of the filter. A recent graph-neural-network study in Cascadia showed that predicting each station only from stations more than 400 km away preserves slow slip, but chose that distance empirically for one region. Here we propose a signal-loss audit of network filtering. We plant transients of known amplitude, width and duration in the real noise of 1,251 continuous stations in and around California, and measure for each filter the fraction of each transient it retains, the spurious signal it creates at nearby stations, and the noise it removes. We audit three families under one leakage-controlled protocol: classical filters that use the target station itself (stacking, robust median imaging and regional principal component analysis), filters that exclude stations within a fixed distance (the 400 km design), and leave-one-out reconstruction models with an exclusion radius varied systematically, from weighted means to a graph-attention network. Fault-consistent transients transplanted from published source models of California events, and the 2019 Ridgecrest postseismic signal, test the conclusions on realistic signals. First results with simple filters show that using the nearest stations removes about 14% of short-period scatter but keeps only about a quarter of a 25-km-wide transient, and that excluding stations within 50 km keeps most of it at little cost in noise removal. The audit will provide trade-off curves and practical guidance for choosing the reach of a filter, and will establish whether deep networks improve the trade-off or only the scatter.

1 Introduction

Continuous GNSS networks record the positions of thousands of geodetic monuments every day, with formal uncertainties of about 1 mm in the horizontal and 3–5 mm in the vertical. The Nevada Geodetic Laboratory (NGL) processes more than 17,000 such stations worldwide with a single, uniform strategy and distributes the time series openly (Blewitt et al., 2018; Nevada Geodetic Laboratory, 2026). These records contain steady interseismic motion, coseismic offsets, postseismic relaxation, slow slip, magmatic inflation and hydrological loading, together with artefacts from antenna and receiver changes. Separating these contributions is the central task of geodetic time-series analysis.

The main obstacle is noise that is shared across stations. Errors in satellite orbits and clocks, in the realisation of the reference frame and in large-scale environmental loading produce a common-mode component of several millimetres that is coherent over hundreds to thousands of kilometres (Wdowinski et al., 1997; Dong et al., 2006; Kreemer and Blewitt, 2021). A wide

range of filters estimates and subtracts this component, from regional stacking (Wdowinski et al., 1997) through principal and independent component analysis (Dong et al., 2006; Shen et al., 2014; Gruszczynski et al., 2018; Zhou et al., 2022; Liu et al., 2026) to correlation-weighted and robust median filters (Tian and Shen, 2016; Kreemer and Blewitt, 2021). These filters are almost always judged by how much they reduce residual scatter or station-to-station correlation, or by how much they shrink velocity uncertainties. None of these criteria distinguishes noise from signal: a filter that also removes a real transient scores better, not worse.

This matters because the transients of interest are themselves spatially coherent. A slow slip event, afterslip or a hydrological load moves many neighbouring stations together, and a filter that estimates the common mode from those neighbours absorbs it. Kreemer and Blewitt (2021) note that much of the scatter reduction of their method comes from local common-mode transients being absorbed into the estimated common mode. The scientific consequences are measurable: in Cascadia, slow-slip scaling laws, b-values and stress drops all change with the detection threshold applied after denoising (Costantino et al., 2026). The disagreement over a reported precursory acceleration before large earthquakes (Bletery and Nocquet, 2023) turned on the same question, whether a common-mode correction removes only noise or also real deformation (Bradley and Hubbard, 2023; Bletery and Nocquet, 2025).

The most direct attempt to control this loss is that of Bachelot et al. (2025). They train a graph neural network to predict each station from other stations, with the target station masked, and subtract the prediction as noise. The decisive design choice is that each station is connected only to its eight nearest stations more than 400 km away, so that the network sees the common mode but not a slow slip event under the target. The 400 km was chosen empirically, as the shortest distance at which 280 slow-slip offsets from 15 Cascadia events were no longer reduced by denoising. A simple average of the stations beyond 400 km performed about as well as the network, and better on one metric, and small offsets were slightly inflated by denoising while large ones were slightly reduced. Their work shows that excluding nearby stations is the key control on signal loss. It leaves open how the right distance depends on the size of the transient, what any filter actually loses when there is no ground truth for the offsets it is compared against, and how the result transfers to other networks and tectonic settings.

Our own preliminary analyses point to the same question. At eleven continuous stations within 80 km of the 2019 Ridgecrest sequence, every apparent anomaly before the mainshock was explained by network-wide noise, an exceptionally wet winter or equipment changes rather than by deformation on the fault. We then built a leakage-controlled benchmark on 1,251 California stations and ran the first planted-signal tests with two simple filters, a distance-weighted mean and a robust weighted median of neighbouring stations. Using the nearest stations, the weighted mean removes about 14% of short-period scatter but keeps only 27% of a 25-km-wide planted transient, 5% of a 50-km-wide one and 36% of the observed Ridgecrest postseismic displacement. Excluding stations within 50 km keeps 83% of the 25-km-wide transient and 74% of the Ridgecrest signal, while still removing 13% of the scatter. Near the rupture, the filter also created spurious transients: at one station 69 km from the epicentre, the cleaned series showed an opposite-sign signal 2.4 times larger than the real one, because its neighbours closer to the rupture had moved more. The robust median kept the same amount of signal as the mean and removed less noise. These numbers were computed with idealised transients on the test years of the benchmark and will be recomputed under the final protocol (Section 3.6), but they show that the trade-off is large and measurable.

Here we propose a signal-loss audit of network filtering. We test the hypothesis that the signal a filter removes is controlled mainly by the ratio of the transient’s footprint to the filter’s reach, and that this holds across filter families, from stacking to deep networks. The study has four objectives. First, we measure, with known transients planted in real noise, how much signal each filter retains and how much spurious signal it creates, as a function of transient width, duration and amplitude. Second, we audit the classical filters in common use and the fixed-distance

design of Bachelot et al. (2025) under the same protocol as leave-one-out reconstruction models with a systematically varied exclusion radius. Third, we test whether deep networks improve the trade-off between noise removed and signal kept, rather than scatter reduction alone. Fourth, we confirm the conclusions on realistic, fault-consistent transients and on the Ridgecrest postseismic signal, and translate them into guidance for choosing the reach of a filter.

2 Literature Review

2.1 Classical common-mode filtering

Network filtering of GNSS data began with the Southern California Integrated GPS Network, where the common-mode error was estimated as the weighted mean of the residuals at the other stations and subtracted from each series (Wdowinski et al., 1997). Principal component analysis and the Karhunen–Loève expansion extended this to common modes that vary in space (Dong et al., 2006). Later work made principal component filtering applicable to real networks, in which stations start and stop at different times: a modified analysis that reproduces each series from its principal components without filling gaps (Shen et al., 2014), probabilistic principal component analysis that treats missing days as latent values (Gruszczynski et al., 2018), and a variational Bayesian version that also selects the number of components, tested on 44 stations in southern California (Li et al., 2020). Because the common mode varies in space, partitioning large networks into sub-regions improves on a single global analysis; on 208 Chinese stations, adaptive sub-regional filtering reduced residual scatter by 20–29%, compared with 15–18% for global principal component analysis (Liu et al., 2026). Independent component analysis offers an alternative decomposition (Gualandi et al., 2016; Zhou et al., 2022).

Tian and Shen (2016) weighted each neighbouring station by its distance and by the correlation of its residuals with the target, chose the weights by a grid search that minimises the residual scatter after filtering, and outperformed principal component analysis by more than 13% in scatter reduction on about 1,600 stations in western North America. They also stepped through minimum-distance thresholds to separate long-range common modes from regional ones and to estimate the spatial extent of regional transients. The common-mode component imaging of Kreemer and Blewitt (2021) takes a correlation-weighted median over nearby stations, including the target station itself when it qualifies as a filter station, and reduced residual scatter by 39–50% on about 2,500 stations in western Europe. Across this family, performance is reported as scatter reduction, reduced inter-station correlation, velocity uncertainty or the recovery of a simulated common mode; none of these measures what the filter removes from real transients.

2.2 Distance-excluded and learned filters

Bachelot et al. (2025) are the closest precedent for this study. Their graph neural network represents stations as nodes and connects each to its eight nearest stations beyond 400 km. During training, 30% of the nodes are masked at random and the mean squared error is computed on the masked nodes only, so the network learns to predict a station from distant stations; the prediction is subtracted as noise. They compare it with three stack filters: the stack of Wdowinski et al. (1997), a stack that excludes stations within 50 km of tremor, and a stack of all stations beyond 400 km. On NGL data the network reduced horizontal common-mode noise by more than 70%, and the stacks performed comparably, the 400 km stack better on the network-average metric. Training used a random 80/20 split of time windows for validation, with 2022–2023 held out for testing.

Costantino et al. (2024) take a supervised route. Their network combines graph recurrent layers and transformers, learns its own adjacency between stations, and is trained to extract slow slip signals from synthetic examples: slip on the subduction interface added to noise synthesised from real Cascadia series by principal component decomposition and Fourier phase

randomisation. Validated by synthetic recovery and by agreement with tremor, it was later used to invert fifteen years of denoised data for the slip history of the interface (Costantino et al., 2026). Because it is taught what a slow slip event looks like, it targets signals of that form; a self-supervised filter that only predicts each station from others makes no assumption about the signal. Mastella et al. (2025) model high-frequency scatter at single stations, with noise defined as the residual of a greedy trajectory decomposition (Bedford and Bevis, 2018) and synthetic training noise generated by an adversarial network trained on about 5,000 NGL stations; spatial correlation is deliberately not used. Teutschmann et al. (2026) pretrained a self-supervised foundation model on 17,651 NGL stations by masked latent prediction, also station by station, and fine-tuned it for displacement forecasting and the localisation of seismic offsets.

Predicting a withheld sensor from its neighbours is a familiar principle elsewhere. Tilt and compliance noise is removed from ocean-bottom seismometer records by predicting the vertical channel from the horizontal and pressure channels (Crawford and Webb, 2000; Janiszewski et al., 2019), distributed acoustic sensing records are denoised by predicting each channel from its neighbours (Batson and Royer, 2019; van den Ende et al., 2023), and graph neural networks for spatiotemporal kriging and imputation are trained by reconstructing randomly withheld sensors (Wu et al., 2021; Cini et al., 2022; Marisca et al., 2022). Excluding nearby observations to prevent spatial autocorrelation from inflating apparent skill is likewise established in buffered spatial cross-validation (Le Rest et al., 2014).

2.3 How signal preservation has been assessed

Table 1 summarises how the studies above test whether a filter preserves real signal. Three approaches recur. Observed transients can be compared before and after filtering, as Bachelot et al. (2025) do with slow-slip offsets; this uses real signals but no ground truth, since the reference offsets are themselves noisy and the comparison cannot say how much of a signal was removed. Synthetic signals can be added to synthetic noise, as in Costantino et al. (2024) and Mastella et al. (2025); this provides ground truth but depends on how realistic the synthetic noise is, and has been analysed as a function of signal-to-noise ratio rather than of spatial scale. Independent data such as tremor can confirm that extracted signals are real (Costantino et al., 2024), but cannot measure what was lost. The classical filters are evaluated on noise alone.

2.4 Transient detection with synthetic training data

Labelled geodetic transients are scarce, so the most successful detectors are trained on synthetic signals. Costantino et al. (2023) injected 60,000 synthetic slow slip events, computed with elastic dislocation models (Okada, 1992), into surrogate noise derived from real Cascadia series, and trained a convolutional, attention-based detector on 135 stations. The detector recovered 35 of 40 events in an independent catalogue (Michel et al., 2019) and reached a detection threshold of about M_w 6.2–6.4. Other strategies include training on real labelled windows, which reached about 75% accuracy in southwest Japan (Tanaka et al., 2025), and convolutional autoencoders for InSAR time series (Rouet-Leduc et al., 2021). Classical network-based detectors include the network strain filter (Ohtani et al., 2010), Kalman smoothing with principal component analysis (Ji and Herring, 2013) and sparse estimation (Riel et al., 2014), and the SCEC exercise compared detection methods on synthetic and real southern California data (Lohman and Murray, 2013). In California, the only systematic statewide analysis of transient deformation is the classical study of Klein et al. (2019), which used about 1,000 stations over 1999–2018 and separated postseismic, hydrological, anthropogenic and magmatic signals. Well-documented California transients suitable as templates include the propagating shallow slow slip events of 2023 on the Superstition Hills and Imperial faults (Materna et al., 2024) and the afterslip of the 2010 El Mayor–Cucapah earthquake (Gonzalez-Ortega et al., 2014).

Table 1: How signal preservation has been assessed in common-mode filtering and GNSS denoising.

Study	Test of signal preservation	Ground truth	Dependence on spatial scale
Tian and Shen (2016)	Minimum residual scatter; scatter against distance threshold	None	Transient extent inferred, signal loss not measured
Kreemer and Blewitt (2021)	Scatter reduction, velocity repeatability	None	Absorption of local transients noted, not measured
PCA and ICA filters ^a	Simulated common mode, inter-station correlation, velocity uncertainty	Noise only	Not examined
Mastella et al. (2025)	Synthetic trajectories in synthetic noise	Yes (synthetic)	Single station
Costantino et al. (2024)	Synthetic slow slip in synthetic noise; tremor	Yes (synthetic)	Analysed against signal-to-noise ratio
Bachelot et al. (2025)	280 observed slow-slip offsets before and after	None	One distance chosen (400 km)
This study	Known transients planted in real noise; fault-consistent transplanted events; Ridgecrest postseismic	Yes (planted)	Retention and spurious signal mapped against footprint and exclusion radius

^a Shen et al. (2014); Gruszczynski et al. (2018); Li et al. (2020); Zhou et al. (2022); Liu et al. (2026).

2.5 Precursory deformation and forecasting

The geodetic evidence for precursory deformation remains contested. The hour-scale acceleration reported by Bletery and Nocquet (2023) from 3,026 high-rate series before 90 earthquakes of $M_w \geq 7$ vanishes after common-mode correction (Bradley and Hubbard, 2023), which the original authors attribute to the correction removing tectonic signal (Bletery and Nocquet, 2025). Stacked tiltmeter records before the 2011 Tohoku-oki earthquake show no precursory signal above a detection limit of about M_w 6.4 (Hirose et al., 2024), and the months-long motions reported before great subduction earthquakes (Bedford et al., 2020) have also been questioned. The best-supported cases are local and slow, such as the eight-month slow slip event before the 2014 Iquique earthquake (Socquet et al., 2017). In earthquake forecasting, geodesy has improved skill mainly through static strain-rate maps combined with smoothed seismicity (Strader et al., 2018; Bayona et al., 2020); afterslip moment does not correlate with aftershock productivity across 41 earthquakes (Churchill et al., 2022), neural forecasting models have not outperformed the epidemic-type aftershock sequence model in California (Stockman et al., 2024), and simple models often match complex networks (Mignan and Broccardo, 2019, 2020). Each of these questions depends on knowing what a filter does to transient signals.

2.6 Synthesis

Five gaps emerge. First, no study has measured with ground truth how much of a transient a common-mode filter removes as a function of the transient’s spatial scale relative to the filter’s reach. Second, the reach that controls signal loss has been chosen empirically, as a fixed 400 km in Cascadia (Bachelot et al., 2025) or by scatter-based thresholds (Tian and Shen, 2016). Third, filters that use the target station, which include most classical methods, have not been compared with leave-one-out filters on signal loss. Fourth, the biases that filtering introduces have been noted only qualitatively, as the inflation of small slow-slip offsets and reduction of large ones after denoising (Bachelot et al., 2025), and spurious signals created by filtering near large transients have not been quantified. Fifth, deep networks have been compared with simple filters on noise

reduction, not at equal signal retention, and network-aware deep denoisers have been developed only for subduction zones.

3 Proposed Method

3.1 Data

We use NGL daily IGS20 solutions for all stations between 32.4° and 42.1°N and between 124.6° and 114.0°W with at least 1,095 daily solutions, which gives 1,381 stations in California and adjacent parts of Nevada, Oregon and Baja California (Nevada Geodetic Laboratory, 2026). We restrict the analysis to 2008 onward, when network density is approximately constant, and remove 130 stations: 30 with less than one year of data after 2008; 22 with non-tectonic site motion such as groundwater-driven subsidence in the Central Valley, identified by a robust residual scale above 10 mm in the horizontal or 30 mm in the vertical; and 78 co-located duplicates within 1 km of a longer-running station. The remaining 1,251 stations provide 6.45×10^6 station-days. In 2018 the median distance to the nearest station was 9.7 km, and a typical station had 23 neighbours within 50 km and 75 within 100 km.

3.2 Preprocessing and data partitioning

For each station i and component $c \in \{E, N, U\}$, we fit by least squares a trajectory model with a constant, a linear rate, annual and semi-annual harmonics and Heaviside steps at catalogued epochs,

$$x_{i,c}(t) = a + v(t - t_0) + \sum_{k=1}^2 [s_k \sin(2\pi kt) + c_k \cos(2\pi kt)] + \sum_j b_j H(t - t_j) + r_{i,c}(t), \quad (1)$$

on the training period only. We subtract the estimated steps at equipment changes but keep earthquake offsets in the residuals, because they are spatially coherent signal. Observations more than five robust standard deviations from a running median are treated as missing. The residuals $r_{i,c}(t)$, in millimetres, with their formal uncertainties and an availability mask, form an array of stations by days by components. We assign 2008–2019 to training and 2020–2021 to validation; 173 stations, spread across the network, are withheld from all training. The interval from July 2019 to June 2020 is excluded from training within 150 km of the Ridgecrest rupture, so that its postseismic signal is unseen by every model.

3.3 Signal-loss accounting

The audit rests on four measurements, computed identically for every filter on cleaned series $r - \hat{r}$, where $\hat{r}$ is the filter’s estimate of the common mode.

Retention. We plant a transient s of known form in the real residuals and filter both the original and the planted data. With $\Delta = (r + s - \hat{r}_{r+s}) - (r - \hat{r}_r)$ the change in the cleaned series caused by the plant, the retention

$$\rho = \frac{\langle \Delta, s \rangle}{\langle s, s \rangle} \quad (2)$$

is the fraction of the transient left in the cleaned data, computed over the horizontal components, the transient’s days and the stations within twice its width. Pairing the planted and unplanted runs removes the noise term that an unpaired comparison would carry; for a linear filter ρ reduces exactly to $\langle s - \hat{s}, s \rangle / \langle s, s \rangle$. The first set of transients has Gaussian footprints, $s_i(t) = A \exp(-d_i^2/2L^2) g(t) \hat{u}$, with amplitudes A of 2–10 mm, widths L of 10–100 km and durations of 10–90 days: a raised-cosine rise, a 60-day hold and a symmetric fall. Fifty transients per

combination, 1,800 in total, are placed at random stations and times and packed into shared runs, separated in space and time so that no station near one transient uses a neighbour near another.

Spurious signal. At stations beyond the footprint of a planted transient whose neighbours lie within it, the planted signal is zero, so any change Δ is created by the filter. We report its peak amplitude as a fraction of A . This quantifies the opposite-sign artefacts we observed near Ridgecrest.

Amplitude dependence. For nonlinear filters, retention can depend on amplitude. Reporting ρ for each A tests the observation of Bachelot et al. (2025) that small offsets were inflated and large ones reduced by denoising.

Noise removed. On the test period, we compute for each station the reduction in short-period scatter, $1 - \text{rms}(\text{fast}(r - \hat{r}))/\text{rms}(\text{fast}(r))$, where the fast part is the series minus its 61-day running median, and take the median over stations.

Plotting noise removed against retention, for each filter and each exclusion radius, gives the trade-off curve that is the central result of the audit. Two further tests use realistic signals. In the event library, published source models of California transients, including the 2023 shallow slow slip events on the Superstition Hills and Imperial faults (Materna et al., 2024), Ridgecrest afterslip and El Mayor–Cucapah afterslip (Gonzalez-Ortega et al., 2014), are transplanted onto other faults of the same type in the SCEC Community Fault Model (Plesch et al., 2024), and the displacement each real station would record is computed with triangular dislocations in an elastic half-space (Nikkhoo and Walter, 2015). Unlike the Gaussian footprints, these reproduce the opposite motion on the two sides of a strike-slip fault, against which a neighbour average partly cancels; until they are in place, results from Gaussian transients are stated for spatially uniform transients only. Finally, we measure the fraction of the observed Ridgecrest postseismic displacement retained at stations within 80 km, on days no model was trained on, and apply the before-and-after offset comparison of Bachelot et al. (2025) to the same signal for direct comparison with their metric.

3.4 Filters audited

Table 2 lists the filters, grouped by how they treat the target station and how their reach is set.

The target-inclusive filters C1–C3 represent common practice: each estimates a station’s common mode from data that include the station itself, so a transient local to the station can be absorbed by construction. Because they see the target, they cannot be scored on reconstruction of withheld data; they are scored only on retention, spurious signal and noise removed. D1 and D2 reproduce the design of Bachelot et al. (2025) on the California network. The leave-one-out models M1–M5 never see the target station and exclude all stations closer than R , which we vary from 0 to 400 km, so that $R = 400$ km reproduces the fixed distance of Bachelot et al. (2025) as one point on a continuum. M1 and M2 are the simple filters of the preliminary results. M3 learns a weight for each neighbour, in a per-station variant and a variant pooled over stations with distance and azimuth as features, so that it also applies to stations absent from training. M4 lets the weights depend on frequency by regressing band-pass filtered series. M5 encodes a 64-day window at each station with a temporal convolutional network and lets the target attend to its neighbours, with attention conditioned on distance and azimuth.

The trainable models are fitted by minimising, over randomly withheld entries $\mathcal{M}$ of the training period,

$$\mathcal{L}(\theta) = \sum_{(i,t,c) \in \mathcal{M}} \frac{(r_{i,c}(t) - \hat{r}_{i,c}(t))^2}{\sigma_{i,c}^2(t) + \sigma_0^2}, \quad (3)$$

with $\sigma_0 = 0.5$ mm. Learned models train on quiet residuals, from which the earthquake offsets estimated in the training-period fit are subtracted; windows within 30 days of catalogued earthquakes are excluded, and residuals are normalised by each station’s noise level. A few stations

Table 2: Filters in the audit.

Label	Filter	Reach	Reference
Target-inclusive (classical)			
C1	Regional stack	Whole region	Wdowinski et al. (1997)
C2	Robust weighted median including the target	Nearest stations	Kreemer and Blewitt (2021)
C3	Sub-regional probabilistic PCA	Sub-region	Gruszczynski et al. (2018); Liu et al. (2026)
Fixed distance exclusion			
D1	Mean of stations beyond 400 km	$d > 400$ km	Bachelot et al. (2025)
D2	Graph neural network, eight nearest beyond 400 km (released code, retrained)	$d > 400$ km	Bachelot et al. (2025)
Leave-one-out with exclusion radius R (this study)			
M1	Distance-weighted mean	$d > R$	Wdowinski et al. (1997)
M2	Robust weighted median	$d > R$	Kreemer and Blewitt (2021)
M3	Ridge regression on neighbours, per station and pooled	$d > R$	Linear spatial prediction
M4	Band-split ridge regression	$d > R$	Multichannel Wiener filter
M5	Graph-attention temporal network	$d > R$	Bachelot et al. (2025); Costantino et al. (2024)

with large coseismic offsets therefore cannot dominate the fit, and the models learn the noise rather than the coherent signal they should leave alone. The same objective drives every learned filter to predict, and therefore remove, any signal its inputs share with the target; this is why the audit measures retention directly rather than relying on the loss. A model that sees the target’s own history can also extrapolate a transient’s onset into withheld days, so cleaned series and all retention scores are computed with the target withheld for the whole period, and M5 is evaluated with and without its own history.

3.5 Comparison with Bachelot et al. (2025)

Table 3 sets out how the proposed audit relates to the closest prior study. The two share the central idea of masking the target and excluding nearby stations; they differ in what is measured and how the exclusion distance is chosen.

3.6 Evaluation protocol

Every model is scored per station and summarised by the median over stations, so that a few stations with large earthquake offsets cannot dominate. Reconstruction errors on withheld data are split into a fast part, the error minus its 61-day running median, and a slow part, and both are normalised by the station’s noise level in the training period. Stations that develop site problems after the training period, an increase in day-to-day scatter by more than a factor of three or an unexplained jump, are left out of scoring. All hyper-parameters and design decisions are made on the validation years. Because our first results were computed on 2022–2026 while the scoring rules were still being revised, those years are no longer untouched; we therefore freeze the method at a fixed date and evaluate it once on NGL data recorded after that date as a prospective test. Differences between filters are assessed with a block bootstrap over days and stations, and trainable models are run with several random seeds. We treat tectonic and other aseismic transients as signal and common-mode error and environmental loading as nuisance;

Table 3: The design of Bachelot et al. (2025) and of this study.

Aspect	Bachelot et al. (2025)	This study
Region and setting	Cascadia, subduction slow slip	California, strike-slip plate boundary
Target station	Masked	Withheld (leave-one-out)
Exclusion distance	Fixed at 400 km	Swept from 0 to 400 km; 400 km included
How the distance is chosen	Shortest distance that preserves observed slow-slip offsets	Read from measured trade-off curves for each transient width
Test of signal preservation	280 observed offsets before and after denoising	Known planted transients, fault-consistent transplanted events, Ridgecrest postseismic signal
Filtering artefacts	Offset bias noted (small offsets inflated, large reduced)	Retention against amplitude; spurious signals measured
Filters compared	Graph network and three stacks	Target-inclusive, fixed-distance and leave-one-out families, linear to deep
Validation	Random 80/20 split of windows; 2022–2023 held out	Temporal split; held-out stations; frozen prospective test

daily loading products (Li et al., 2025) allow the loading contribution to the common mode to be estimated separately.

3.7 Extension to continual transient detection

The filter with the best trade-off will provide the residuals for a continual detector: a multi-scale cumulative-sum statistic combined with a spatial-coherence criterion, calibrated to a set number of false alarms per station-year and run in a strictly chronological replay with model updates restricted to windows not flagged as anomalous (Kirkpatrick et al., 2017). This extension depends on the outcome of the audit.

4 Expected Findings

We design the study so that every outcome is informative, and we test four hypotheses.

First, we expect every filter that uses the target station or its nearest neighbours to remove most of any transient wider than the station spacing. In the preliminary results the leave-one-out weighted mean kept about a quarter of a 25-km-wide transient; target-inclusive filters such as stacking, robust median imaging and principal component analysis should keep no more, and less of transients local to a single station, while appearing to remove more noise.

Second, we expect retention to depend mainly on the ratio of transient width to exclusion radius, rising from near zero when the width exceeds the radius to near one when the radius exceeds it, while noise removal declines more slowly because the common mode is coherent over larger distances. The preliminary results follow this pattern. The 400 km design of Bachelot et al. (2025) should then retain transients up to a few hundred kilometres wide but remove less local noise than intermediate radii; for California, where many transients are a few tens of kilometres wide, we expect a radius of about 25–50 km to give a better balance. Fault-consistent transients should be retained somewhat better than Gaussian ones at the same width, because neighbouring stations on opposite sides of a fault move in opposite directions.

Third, we expect filters to create spurious transients near large real signals, of opposite sign and up to the size of the real signal, at small radii, and these to disappear once the radius exceeds the signal’s extent.

Fourth, we expect learned models to reduce reconstruction error relative to the weighted mean, but to improve the trade-off between noise removed and signal kept only modestly, if at all. Consistent with Bachelot et al. (2025), where a simple stack matched a graph network, a deep model that wins only on scatter would be absorbing more signal, and the audit would show it.

The main deliverables are a signal-loss table and trade-off curves for the filter families in common use, measured with ground truth on a dense strike-slip network; a quantitative basis for choosing the reach of a filter for transients of a given size; measured rates of spurious signals; and an open benchmark with data, splits, planted signals and code. Whatever the ranking of the filters, these results will let geodesists report what their filtering costs in signal alongside what it removes in noise.

AI Usage Declaration

An AI assistant, Claude Code running the Claude Opus 5.5 model (Anthropic), was used to help design the benchmark, write the software and its tests, run the preliminary analyses and produce the figures, search and summarise the literature, and draft and edit this text. Bibliographic entries were generated from Crossref and arXiv metadata. The research question and the central idea of predicting withheld stations from their neighbours are the author's. The author directed the work, reviewed all AI-generated material and is responsible for the content. Some cited papers have so far been assessed from their abstracts and will be read in full before any results are submitted for publication.

References

- Loïc Bachelot, Amanda Thomas, Diego Melgar, Jake Searcy, and Yu-Sheng Sun. Cascadia Daily GNSS Time Series Denoising: Graph Neural Network and Stack Filtering. *Seismica*, 4 (1), 2025. doi: 10.26443/seismica.v2i4.1419.
- Joshua Batson and Loïc Royer. Noise2Self: Blind Denoising by Self-Supervision. In *Proceedings of the 36th International Conference on Machine Learning*, volume 97 of *Proceedings of Machine Learning Research*, pages 524–533, 2019.
- J A Bayona, W Savran, A Strader, S Hainzl, F Cotton, and D Schorlemmer. Two global ensemble seismicity models obtained from the combination of interseismic strain measurements and earthquake-catalogue information. *Geophysical Journal International*, 224(3):1945–1955, 2020. doi: 10.1093/gji/ggaa554.
- Jonathan Bedford and Michael Bevis. Greedy Automatic Signal Decomposition and Its Application to Daily GPS Time Series. *Journal of Geophysical Research: Solid Earth*, 123(8): 6992–7003, 2018. doi: 10.1029/2017JB014765.
- Jonathan R. Bedford, Marcos Moreno, Zhiguo Deng, Onno Oncken, Bernd Schurr, Timm John, Juan Carlos Báez, and Michael Bevis. Months-long thousand-kilometre-scale wobbling before great subduction earthquakes. *Nature*, 580(7805):628–635, 2020. doi: 10.1038/s41586-020-2212-1.
- Quentin Bletery and Jean-Mathieu Nocquet. The precursory phase of large earthquakes. *Science*, 381(6655):297–301, 2023. doi: 10.1126/science.adg2565.
- Quentin Bletery and Jean-Mathieu Nocquet. Do large earthquakes start with a precursory phase of slow slip? *Seismica*, 3(2), 2025. doi: 10.26443/seismica.v3i2.1383.

- Geoffrey Blewitt, William Hammond, and Corn  Kreemer. Harnessing the GPS Data Explosion for Interdisciplinary Science. *Eos*, 99, 2018. doi: 10.1029/2018EO104623.
- Kyle Bradley and Judith Hubbard. Earthquake precursors? Not so fast. *Earthquake Insights*, 2023.
- R. M. Churchill, M. J. Werner, J. Biggs, and  . Fagereng. Relative Afterslip Moment Does Not Correlate With Aftershock Productivity: Implications for the Relationship Between Afterslip and Aftershocks. *Geophysical Research Letters*, 49(24):e2022GL101165, 2022. doi: 10.1029/2022GL101165.
- Andrea Cini, Ivan Marisca, and Cesare Alippi. Filling the G_ap_s: Multivariate Time Series Imputation by Graph Neural Networks. In *International Conference on Learning Representations (ICLR)*, 2022. arXiv:2108.00298.
- Giuseppe Costantino, Sophie Giffard-Roisin, Mathilde Radiguet, Mauro Dalla Mura, David Marsan, and Anne Socquet. Multi-station deep learning on geodetic time series detects slow slip events in Cascadia. *Communications Earth & Environment*, 4(1):435, 2023. doi: 10.1038/s43247-023-01107-7.
- Giuseppe Costantino, Sophie Giffard-Roisin, Mauro Dalla Mura, and Anne Socquet. Denoising of Geodetic Time Series Using Spatiotemporal Graph Neural Networks: Application to Slow Slip Event Extraction. *IEEE Journal of Selected Topics in Applied Earth Observations and Remote Sensing*, 17:17567–17579, 2024. doi: 10.1109/JSTARS.2024.3465270.
- Giuseppe Costantino, Mathilde Radiguet, Zaccaria El Yousfi, and Anne Socquet. A Continuum of Slow Slip Events in the Cascadia Subduction Zone Illuminated by High-Resolution Deep-Learning Denoising. *Geophysical Research Letters*, 53(3):e2025GL117446, 2026. doi: 10.1029/2025GL117446.
- W. C. Crawford and S. C. Webb. Identifying and Removing Tilt Noise from Low-Frequency (<0.1 Hz) Seafloor Vertical Seismic Data. *Bulletin of the Seismological Society of America*, 90(4):952–963, 2000. doi: 10.1785/0119990121.
- D. Dong, P. Fang, Y. Bock, F. Webb, L. Prawirodirdjo, S. Kedar, and P. Jamason. Spatiotemporal filtering using principal component analysis and Karhunen-Loeve expansion approaches for regional GPS network analysis. *Journal of Geophysical Research: Solid Earth*, 111(B3):2005JB003806, 2006. doi: 10.1029/2005JB003806.
- Alejandro Gonzalez-Ortega, Yuri Fialko, David Sandwell, F. Alejandro Nava-Pichardo, John Fletcher, Javier Gonzalez-Garcia, Brad Lipovsky, Michael Floyd, and Gareth Funning. El Mayor-Cucapah (Mw7.2) earthquake: Early near-field postseismic deformation from InSAR and GPS observations. *Journal of Geophysical Research: Solid Earth*, 119(2):1482–1497, 2014. doi: 10.1002/2013jb010193.
- Maciej Gruszczynski, Anna Klos, and Janusz Bogusz. A Filtering of Incomplete GNSS Position Time Series with Probabilistic Principal Component Analysis. *Pure and Applied Geophysics*, 175(5):1841–1867, 2018. doi: 10.1007/s00024-018-1856-3.
- A. Gualandi, E. Serpelloni, and M. E. Belardinelli. Blind source separation problem in GPS time series. *Journal of Geodesy*, 90(4):323–341, 2016. doi: 10.1007/s00190-015-0875-4.
- Hitoshi Hirose, Aitaro Kato, and Takeshi Kimura. Did Short-Term Preseismic Crustal Deformation Precede the 2011 Great Tohoku-Oki Earthquake? An Examination of Stacked Tilt Records. *Geophysical Research Letters*, 51(12):e2024GL109384, 2024. doi: 10.1029/2024GL109384.

- Helen A Janiszewski, James B Gaherty, Geoffrey A Abers, Haiying Gao, and Zachary C Eilon. Amphibious surface-wave phase-velocity measurements of the Cascadia subduction zone. *Geophysical Journal International*, 217(3):1929–1948, 2019. doi: 10.1093/gji/ggz051.
- Kang Hyeun Ji and Thomas A. Herring. A method for detecting transient signals in GPS position time-series: smoothing and principal component analysis. *Geophysical Journal International*, 193(1):171–186, 2013. doi: 10.1093/gji/ggt003.
- James Kirkpatrick, Razvan Pascanu, Neil Rabinowitz, Joel Veness, Guillaume Desjardins, Andrei A. Rusu, Kieran Milan, John Quan, Tiago Ramalho, Agnieszka Grabska-Barwinska, Demis Hassabis, Claudia Clopath, Dharshan Kumaran, and Raia Hadsell. Overcoming catastrophic forgetting in neural networks. *Proceedings of the National Academy of Sciences*, 114(13):3521–3526, 2017. doi: 10.1073/pnas.1611835114.
- Emilie Klein, Yehuda Bock, Xiaohua Xu, David T. Sandwell, Dorian Golriz, Peng Fang, and Lina Su. Transient Deformation in California From Two Decades of GPS Displacements: Implications for a Three-Dimensional Kinematic Reference Frame. *Journal of Geophysical Research: Solid Earth*, 124(11):12189–12223, 2019. doi: 10.1029/2018JB017201.
- Corné Kreemer and Geoffrey Blewitt. Robust estimation of spatially varying common-mode components in GPS time-series. *Journal of Geodesy*, 95(1):13, 2021. doi: 10.1007/s00190-020-01466-5.
- Kévin Le Rest, David Pinaud, Pascal Monestiez, Joël Chadoeuf, and Vincent Bretagnolle. Spatial leave-one-out cross-validation for variable selection in the presence of spatial autocorrelation. *Global Ecology and Biogeography*, 23(7):811–820, 2014. doi: 10.1111/geb.12161.
- Wudong Li, Weiping Jiang, Zhao Li, Hua Chen, Qusen Chen, Jian Wang, and Guangbin Zhu. Extracting Common Mode Errors of Regional GNSS Position Time Series in the Presence of Missing Data by Variational Bayesian Principal Component Analysis. *Sensors*, 20(8):2298, 2020. doi: 10.3390/s20082298.
- Zhao Li, Weiping Jiang, Tonie van Dam, Xiaowei Zou, Qusen Chen, and Hua Chen. A Review on Modeling Environmental Loading Effects and Their Contributions to Nonlinear Variations of Global Navigation Satellite System Coordinate Time Series. *Engineering*, 47:26–37, 2025. doi: 10.1016/j.eng.2024.09.001.
- Bin Liu, Yifan Jiang, Zien Xiao, Weilong Rao, and Tengfei Feng. An adaptive spatiotemporal filtering method for GNSS coordinate time series of CMONOC. *Earth, Planets and Space*, 78(1):158, 2026. doi: 10.1186/s40623-026-02477-w.
- R. B. Lohman and J. R. Murray. The SCEC Geodetic Transient-Detection Validation Exercise. *Seismological Research Letters*, 84(3):419–425, 2013. doi: 10.1785/0220130041.
- Ivan Marisca, Andrea Cini, and Cesare Alippi. Learning to Reconstruct Missing Data from Spatiotemporal Graphs with Sparse Observations. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2022. arXiv:2205.13479.
- Giacomo Mastella, Jonathan Bedford, Fabio Corbi, and Francesca Funiciello. Denoising daily displacement GNSS time series using deep neural networks in a near real-time framing: a single-station method. *Geophysical Journal International*, 242(3):ggaf207, 2025. doi: 10.1093/gji/ggaf207.
- Kathryn Materna, Roland Bürgmann, Danielle Lindsay, Roger Bilham, Thomas Herring, Brendan Crowell, and Walter Szeliga. Shallow Slow Slip Events in the Imperial Valley With Along-Strike Propagation. *Geophysical Research Letters*, 51(12):e2023GL108089, 2024. doi: 10.1029/2023gl108089.

- Sylvain Michel, Adriano Gualandi, and Jean-Philippe Avouac. Similar scaling laws for earthquakes and Cascadia slow-slip events. *Nature*, 574(7779):522–526, 2019. doi: 10.1038/s41586-019-1673-6.
- Arnaud Mignan and Marco Broccardo. One neuron versus deep learning in aftershock prediction. *Nature*, 574(7776):E1–E3, 2019. doi: 10.1038/s41586-019-1582-8.
- Arnaud Mignan and Marco Broccardo. Neural Network Applications in Earthquake Prediction (1994–2019): Meta-Analytic and Statistical Insights on Their Limitations. *Seismological Research Letters*, 91(4):2330–2342, 2020. doi: 10.1785/0220200021.
- Nevada Geodetic Laboratory. GPS Networks Map Viewer and Time Series Products (IGS20 daily tenv3). <https://geodesy.unr.edu>, 2026. Accessed 6 October 2026.
- Mehdi Nikkhoo and Thomas R. Walter. Triangular dislocation: an analytical, artefact-free solution. *Geophysical Journal International*, 201(2):1119–1141, 2015. doi: 10.1093/gji/ggv035.
- Ryu Ohtani, Jeffrey J. McGuire, and Paul Segall. Network strain filter: A new tool for monitoring and detecting transient deformation signals in GPS arrays. *Journal of Geophysical Research: Solid Earth*, 115(B12), 2010. doi: 10.1029/2010JB007442.
- Yoshimitsu Okada. Internal deformation due to shear and tensile faults in a half-space. *Bulletin of the Seismological Society of America*, 82(2):1018–1040, 1992. doi: 10.1785/BSSA0820021018.
- Andreas Plesch, Scott Marshall, and John Shaw. SCEC Community Fault Model (CFM), version 7.0. Zenodo, 2024.
- Bryan Riel, Mark Simons, Piyush Agram, and Zhongwen Zhan. Detecting transient signals in geodetic time series using sparse estimation techniques. *Journal of Geophysical Research: Solid Earth*, 119(6):5140–5160, 2014. doi: 10.1002/2014JB011077.
- Bertrand Rouet-Leduc, Romain Jolivet, Manon Dalaison, Paul A. Johnson, and Claudia Hubbert. Autonomous extraction of millimeter-scale deformation in InSAR time series using deep learning. *Nature Communications*, 12(1):6480, 2021. doi: 10.1038/s41467-021-26254-3.
- Yunzhong Shen, Weiwei Li, Guochang Xu, and Bofeng Li. Spatiotemporal filtering of regional GNSS network’s position time series with missing data using principle component analysis. *Journal of Geodesy*, 88(1):1–12, 2014. doi: 10.1007/s00190-013-0663-y.
- Anne Socquet, Jesus Piña Valdes, Jorge Jara, Fabrice Cotton, Andrea Walpersdorf, Nathalie Cotte, Sebastian Specht, Francisco Ortega-Culaciati, Daniel Carrizo, and Edmundo Norabuena. An 8 month slow slip event triggers progressive nucleation of the 2014 Chile megathrust. *Geophysical Research Letters*, 44(9):4046–4053, 2017. doi: 10.1002/2017GL073023.
- Samuel Stockman, Daniel Lawson, and Maximilian Werner. EarthquakeNPP: A Benchmark for Earthquake Forecasting with Neural Point Processes. arXiv preprint arXiv:2410.08226, 2024.
- Anne Strader, Maximilian Werner, José Bayona, Philip Maechling, Fabio Silva, Maria Liukis, and Danijel Schorlemmer. Prospective Evaluation of Global Earthquake Forecast Models: 2 Yrs of Observations Provide Preliminary Support for Merging Smoothed Seismicity with Geodetic Strain Rates. *Seismological Research Letters*, 89(4):1262–1271, 2018. doi: 10.1785/0220180051.
- Yusuke Tanaka, Masayuki Kano, and Keisuke Yano. Detecting Slow Slip Signals in Southwest Japan Based on Machine Learning Trained by Real GNSS Time Series. *Journal of Geophysical Research: Solid Earth*, 130(2):e2024JB029499, 2025. doi: 10.1029/2024JB029499.

- Nick Teutschmann, Laura Crocetti, Fanny Lehmann, Leonardo Trentini, and Benedikt Soja. GNSS-DispFM: A Self-Supervised Foundation Model for Daily GNSS Displacement Time Series. *IEEE Journal of Selected Topics in Applied Earth Observations and Remote Sensing*, 2026. doi: 10.1109/JSTARS.2026.3740122. Early access.
- Yunfeng Tian and Zheng-Kang Shen. Extracting the regional common-mode component of GPS station position time series from dense continuous network. *Journal of Geophysical Research: Solid Earth*, 121(2):1080–1096, 2016. doi: 10.1002/2015JB012253.
- Martijn van den Ende, Itzhak Lior, Jean-Paul Ampuero, Anthony Sladen, André Ferrari, and Cédric Richard. A Self-Supervised Deep Learning Approach for Blind Denoising and Waveform Coherence Enhancement in Distributed Acoustic Sensing Data. *IEEE Transactions on Neural Networks and Learning Systems*, 34(7):3371–3384, 2023. doi: 10.1109/tnnls.2021.3132832.
- Shimon Wdowinski, Yehuda Bock, Jie Zhang, Peng Fang, and Joachim Genrich. Southern California permanent GPS geodetic array: Spatial filtering of daily positions for estimating coseismic and postseismic displacements induced by the 1992 Landers earthquake. *Journal of Geophysical Research: Solid Earth*, 102(B8):18057–18070, 1997. doi: 10.1029/97JB01378.
- Yuankai Wu, Dingyi Zhuang, Aurelie Labbe, and Lijun Sun. Inductive Graph Neural Networks for Spatiotemporal Kriging. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 35, pages 4478–4485, 2021. doi: 10.1609/aaai.v35i5.16575.
- Wei Zhou, Kaihua Ding, Peng Liu, Guanghong Lan, and Zutao Ming. Spatiotemporal Filtering for Continuous GPS Coordinate Time Series in Mainland China by Using Independent Component Analysis. *Remote Sensing*, 14(12):2904, 2022. doi: 10.3390/rs14122904.